

KUCC - 머신러닝 & 딥러닝 입문

머신러닝&딥러닝 세션

—
2주차

| 머신러닝 훈련 (1)

참고 사항

| 본 교안은 '핸즈온 머신러닝' 내용을 참고하고, 정리한 것입니다.

| GitHub 코드 링크는 아래와 같으며, 더 많은 공부를 원하시는 분들은 해당 링크에서 다른 내용에 대해 깊게 공부할 수 있습니다.

| Icon 사용 출처 : <https://www.FlatIcon.com/kr>

[핸즈온 머신러닝 깃허브] <https://github.com/ExcelsiorCJH/Hands-On-ML>

[Ratsgo님 깃허브] <https://ratsgo.github.io/blog/>

목차

0

—
머신 러닝의 개념

1

—
데이터 가져오기

2

—
데이터 통찰/시각화


3

—
데이터 훈련/평가

머신 러닝의 개념

1.2 왜 머신러닝을 사용하는가?

- 머신러닝 모델을 통해 코드를 간단하고 더 잘 수행할 수 있도록 함
- 머신러닝 기법으로 복잡한 문제를 해결할 수 있음
- 새로운 데이터에 대해 유동적으로 적응할 수 있음



머신 러닝의 개념


1.3 머신 러닝의 시스템과 종류

- | 사람의 감독 하에 훈련하는 것인가?
 - ▶ 지도/비지도/준 지도/강화 학습
- | 실시간으로 점진적인 학습을 하는 것인지
 - ▶ 온라인/배치 학습
- | 단순한 데이터 비교인지, 패턴을 발견하는 모델을 만드는지
 - ▶ 사례/모델 기반 학습

머신 러닝의 개념


1.3.1 지도 학습과 비지도 학습

- '학습하는 동안의 감독 형태나 정보량'에 따라 머신러닝 종류를 분류



머신 러닝의 개념

〈머신러닝의 종류〉



머신 러닝의 개념

1.3.2 배치 학습과 온라인 학습

- **입력 데이터의 스트림**으로부터 점진적으로 학습할 수 있는지 여부로 분류


배치 학습 (Batch learning) : **정진적으로 학습X**

- 점진적으로 학습할 수 없으며, 데이터를 모두 사용해 훈련 시킴
- 주로 오프라인에서 수행 → **offline learning**


온라인 학습 (Online learning) : **정진적으로 학습**

- 데이터를 순차적 또는 작은 묶음 단위(미니 배치)로 학습
- 빠른 변화에 적응해야 하는 시스템에 적합(ex. 주식가격)

Batch Learning



On-line Learning



머신 러닝의 개념

1.3.3 사례 기반 학습과 모델 기반 학습



- 어떻게 일반화 되는가에 따라 분류
- 일반화란 학습데이터가 아닌 새로운 데이터를 예측하는데 전반적으로 잘 예측할 수 있도록 하는 것을 말함

① 사례 기반 학습 (Instance-based learning)

- 이미 알고 있는 데이터와 새로운 데이터 간의 유사한 정도를 통해 예측하는 방법
- 대표적인 알고리즘으로는 k-NN(k-Nearest Neighbors)가 있음

② 모델 기반 학습 (Model-based learning)

- 학습 데이터로부터 일반화 할 수 있는 모델을 만들어 예측하는 방법



머신 러닝의 개념

1.4 머신 러닝의 주요 도전 과제 : Data

| 머신 러닝 Algorithm이 잘 작동하기 위해서는 Data가 많아야 함

- ▶ 빅데이터 기반의 학습

| 일반화가 잘 되려면 새로운 데이터를 학습데이터가 잘 대표하여야 함

- ▶ 학습데이터의 정규성/보편성/등분산성

| 학습 데이터 정제에 많은 시간을 투자해야 함

- ▶ 학습 데이터에 존재하는 에러(이상치, 영향치, 노이즈) 제거 필요

| 학습에 사용할 좋은 특성을 찾는 것이 중요 (by using Feature Engineering)

- ▶ 특성 선택 (feature selection) : 가지고 있는 특성 중 유용한 특성 선택
- ▶ 특성 추출 (feature extraction) : 특성을 결합하여 유용한 특성 생성

머신 러닝의 개념


1.4 머신러닝의 주요 도전 과제 - 알고리즘

↪ train set에 너무 빠져함

- 머신러닝 알고리즘이 학습데이터에 **과대적합(overfitting)** 되지 않아야 한다.
 - 파라미터 수가 적은 모델을 선택(고차원 다항 모델 보다 선형모델)
 - 학습 데이터에서 특성 수를 줄이거나, 모델에 제약을 가함
 - 학습 데이터를 더 많이 모은다.
 - 학습 데이터의 노이즈를 줄인다.

↪ test set에 너무 빠져함

- 또한, **과소적합(underfitting)** 되지 않아야 한다.
 - 파라미터가 더 많은 모델을 선택
 - 더 좋은 특성을 추가
 - 모델의 제약을 줄임



머신 러닝의 개념


1.5 정리

- | 머신 러닝은 명시적인 규칙을 코딩하는 것이 아니다. 기계가 데이터로부터 학습하여, 특정 작업을 잘하도록 만드는 것이다.
- | 머신 러닝은 지도/비지도 학습, 배치/온라인 학습, 사례/모델 기반 학습으로 분류 가능하다.
- | 머신 러닝은 훈련(학습) 세트에 데이터를 모아 학습 알고리즘에 넣어준다.
 - ▶ 모델 기반 : 훈련 세트에 모델을 맞추기 위해 파라미터 조정
 - ▶ 사례 기반 : 학습 데이터를 기억하고, 새로운 데이터를 일반화하기 위해 유사도 측정
- | 훈련 세트가 많고, 훈련 세트가 새로운 데이터를 대표할 수 있으며, 노이즈가 별로 없는 데이터로 학습을 진행해야 함.

머신 러닝의 개념

1.6 테스트와 검증

- **학습 세트(Train Set):** 머신러닝 모델을 학습할 때 사용하는 데이터 셋 **— 학습**
- **검증 세트(Validation Set):** 모델학습에 필요한 하이퍼파라미터를 찾기위해 사용하는 데이터 셋 **— 하이퍼파라미터 찾기**
- **테스트 세트(Test Set):** 학습된 모델을 평가하기 위한 데이터 셋 **— 평가**



캘리포니아 주택 가격 예측!



1.5 정리

| 블록 그룹마다 중간소득/인구/중간 주택 가격 有

| 파이프라인 : 데이터 처리 컴포넌트의 연속

- 각 컴포넌트는 데이터를 추출해 목적에 맞게 처리
- 결과를 다른 컴포넌트나 저장소에 보내는 일련의 과정

1

데이터 가져오기(해킹?)

기본 Library Import

라이브러리 설명

```
In [1]: import pandas as pd
import numpy as np
import os

np.random.seed(42)

import matplotlib.pyplot as plt
import platform

%matplotlib inline
```

numpy

- ▶ 행렬(Matrix) 연산 관련 라이브러리
- ▶ 선형대수 구현 라이브러리!

pandas

- ▶ 데이터프레임(Table) 관련 라이브러리

matplotlib

- ▶ 시각화 관련 라이브러리(그래프)

os, platform...

- ▶ 운영체제, 시스템 정보 관련 라이브러리

기본 Library Import

시각화 위해 한글적용코드

```
path = 'c:/windows/Fonts/malgun.ttf'

from matplotlib import font_manager, rc

if platform.system() == 'Darwin':          ) mac 폰트
    rc('font', family = 'AppleGothic')
elif platform.system() == 'Windows':        ) 윈도우 폰트
    font_name = font_manager.FontProperties(fname=path).get_name()
    rc('font', family=font_name)
else:
    print('unknown system...sorry~~~') — mac/윈도우가 아닌 경우

plt.rcParams['axes.unicode_minus'] = False
```


한글 적용 코드

Matplotlib에서 한글을 사용하려면 폰트 지정을 해줘야 합니다.

데이터 가져오기

ExcelsiorCJH DQN - Keras		
	Author	Date
📁 Chap01-The_Machine_Learning_Land...	cheat sheet add	3 years ago
📁 Chap02-End_to_End_ML_Project	Chap02 - 원료	3 years ago
📁 Chap03-Classification	Chap03 - Classification	3 years ago
📁 Chap04-Training_Models	Chap04 - Model Training	3 years ago
📁 Chap05-SVM	오타수정	3 years ago
📁 Chap06-Decision_Tree	오타 수정	3 years ago
📁 Chap07-Ensemble_Learning_and_Ran...	markdown add	3 years ago
📁 Chap08-Dimensionality_Reduction	재실행	3 years ago
📁 Chap09-Up_and_Running_with_Tenso...	Running in Mac	3 years ago
📁 Chap10-Introduction_to_ANN	chap10 - img & markdown	3 years ago
📁 Chap11-Training_DNN	오타 수정	3 years ago
📁 Chap13-Convolutional_Neural_Netwo...	발표	2 years ago
📁 Chap14-Recurrent_Neural_Networks	오타 수정	2 years ago
📁 Chap15-Autoencoders	설명 추가	2 years ago
📁 Chap16-Reinforcement_Learning	DQN - Keras	2 years ago
📁 datasets	dataset 추가	3 years ago
📄 .gitignore	python add	2 years ago
📄 README.md	Chap16 - RL	2 years ago
📄 cover.PNG	cover & README.md 추가	3 years ago

데이터 가져오기



master Hands-On-ML / datasets / housing / Go to file

ExcelsiorCJH dataset 추가 c63a6dd on 26 Jul 2018 History

..

README.md dataset 추가 3 years ago

housing.csv dataset 추가 3 years ago housing.csv

housing.tgz dataset 추가 3 years ago

California Housing

Source

This dataset is a modified version of the California Housing dataset available from [Luis Torgo's page](#) (University of Porto). Luis Torgo obtained it from the StatLib repository (which is closed now). The dataset may also be downloaded from StatLib mirrors.

This dataset appeared in a 1997 paper titled *Sparse Spatial Autoregressions* by Pace, R. Kelley and Ronald Barry, published in the *Statistics and Probability Letters* journal. They built it using the 1990 California census data. It contains one row per census block group. A block group is the smallest geographical unit for which the U.S. Census Bureau publishes sample data (a block group typically has a population of 600 to 3,000 people).

Tweaks

The dataset in this directory is almost identical to the original, with two differences:

- 207 values were randomly removed from the `total_bedrooms` column, so we can discuss what to do with missing data.
- An additional categorical attribute called `ocean_proximity` was added, indicating (very roughly) whether each block group is near the ocean, near the Bay area, inland or on an island. This allows discussing what to do with categorical data.

Note that the block groups are called "districts" in the Jupyter notebooks, simply because in some contexts the name "block group" was confusing.

Data description

```
>>> housing.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
longitude      20640 non-null float64
latitude       20640 non-null float64
housing_median_age 20640 non-null float64
total_rooms    20640 non-null float64
total_bedrooms 20433 non-null float64
population     20640 non-null float64
households     20640 non-null float64
median_income   20640 non-null float64
median_house_value 20640 non-null float64
ocean_proximity 20640 non-null object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB

>>> housing["ocean_proximity"].value_counts()
<1H OCEAN    9136
INLAND       6551
NEAR OCEAN    2658
NEAR BAY      2290
ISLAND        5
Name: ocean_proximity, dtype: int64

>>> housing.describe()
   longitude      latitude  housing_median_age  total_rooms \
count  16513.000000  16513.000000  16513.000000  16513.000000
mean   -119.575972   35.639693   28.652335  2622.347605
std     2.002048    2.138279   12.576306  2138.559393
min    -124.350000   32.540000   1.000000   6.000000
25%    -121.800000   33.940000   18.000000  1442.000000
50%    -118.510000   34.260000   29.000000  2119.000000
75%    -118.010000   37.720000   37.000000  3141.000000
max    -114.310000   41.950000   52.000000  39320.000000


   total_bedrooms  population  households  median_income
count  16355.000000  16513.000000  16513.000000  16513.000000
mean    534.885112   1419.525465   496.975050   3.875651
std     412.716467   1115.715084   375.737945   1.905088
min     2.000000     3.000000     2.000000   0.499900
25%    295.000000    784.000000   278.000000   2.566800
50%    433.000000   1164.000000   408.000000   3.541400
75%    644.000000   1718.000000   602.000000   4.745000
max    6210.000000  35682.000000  5358.000000  15.000100
```

데이터 가져오기

longitude,latitude,housing_median_age,total_rooms,total_bedrooms,population,households,median_income,median_house_value,ocean_proximity

-122.23,37.88,41.0,880.0,129.0,322.0,126.0,8.3252,452600.0,NEAR BAY
-122.22,37.86,21.0,7099.0,1106.0,2401.0,1138.0,8.3014,358500.0,NEAR BAY
-122.24,37.85,52.0,1467.0,190.0,496.0,177.0,7.2574,352100.0,NEAR BAY
-122.25,37.85,52.0,1274.0,235.0,558.0,219.0,5.6431,341300.0,NEAR BAY
-122.25,37.85,52.0,1627.0,280.0,556.0,259.0,3.8462,342200.0,NEAR BAY
-122.25,37.85,52.0,919.0,213.0,413.0,193.0,4.0368,269700.0,NEAR BAY
-122.25,37.84,52.0,2535.0,489.0,1094.0,514.0,3.6591,299200.0,NEAR BAY
-122.25,37.84,52.0,3104.0,687.0,1157.0,647.0,3.12,241400.0,NEAR BAY
-122.26,37.84,42.0,2555.0,665.0,1206.0,595.0,2.0804,226700.0,NEAR BAY
-122.25,37.84,52.0,3549.0,707.0,1551.0,714.0,3.6912,261100.0,NEAR BAY
-122.26,37.85,52.0,2202.0,434.0,910.0,402.0,3.2031,281500.0,NEAR BAY
-122.26,37.85,52.0,3503.0,752.0,1504.0,734.0,3.2705,241800.0,NEAR BAY
-122.26,37.85,52.0,2491.0,474.0,1098.0,468.0,3.075,213500.0,NEAR BAY
-122.26,37.84,52.0,696.0,191.0,345.0,174.0,2.6736,191300.0,NEAR BAY
-122.26,37.85,52.0,2643.0,626.0,1212.0,620.0,1.9167,159200.0,NEAR BAY
-122.26,37.85,50.0,1120.0,283.0,697.0,264.0,2.125,140000.0,NEAR BAY
-122.27,37.85,52.0,1966.0,347.0,793.0,331.0,2.775,152500.0,NEAR BAY
-122.27,37.85,52.0,1228.0,293.0,648.0,303.0,2.1202,155500.0,NEAR BAY
-122.26,37.84,50.0,2239.0,455.0,990.0,949.0,1.9911,158700.0,NEAR BAY
-122.27,37.84,52.0,1503.0,298.0,690.0,275.0,2.6033,162900.0,NEAR BAY
-122.27,37.85,40.0,751.0,184.0,409.0,166.0,1.3578,147500.0,NEAR BAY
-122.27,37.85,42.0,1639.0,367.0,929.0,366.0,1.7135,159800.0,NEAR BAY
-122.27,37.84,52.0,2436.0,541.0,1015.0,478.0,1.725,113900.0,NEAR BAY
-122.27,37.84,52.0,1688.0,337.0,853.0,325.0,2.1806,99700.0,NEAR BAY
-122.27,37.84,52.0,2224.0,437.0,1006.0,422.0,2.6,132600.0,NEAR BAY
-122.28,37.85,41.0,535.0,123.0,317.0,119.0,2.4038,107500.0,NEAR BAY
-122.28,37.85,49.0,1130.0,244.0,607.0,239.0,2.4597,93800.0,NEAR BAY
-122.28,37.85,52.0,1898.0,421.0,1102.0,397.0,1.808,105500.0,NEAR BAY
-122.28,37.84,50.0,2082.0,492.0,1131.0,473.0,1.6424,108900.0,NEAR BAY
-122.28,37.84,52.0,729.0,160.0,395.0,155.0,1.6875,132000.0,NEAR BAY
-122.28,37.84,49.0,1916.0,447.0,863.0,378.0,1.9274,122300.0,NEAR BAY
-122.28,37.84,52.0,2153.0,481.0,1168.0,441.0,1.9615,115200.0,NEAR BAY
-122.27,37.84,48.0,1922.0,409.0,1026.0,335.0,1.7969,110400.0,NEAR BAY
-122.27,37.83,49.0,1655.0,366.0,754.0,329.0,1.375,104900.0,NEAR BAY
-122.27,37.83,51.0,2665.0,574.0,1258.0,536.0,2.7303,109700.0,NEAR BAY
-122.27,37.83,49.0,1215.0,282.0,570.0,264.0,1.4861,97200.0,NEAR BAY
-122.27,37.83,48.0,1798.0,432.0,987.0,374.0,1.0972,104500.0,NEAR BAY
-122.28,37.83,52.0,1511.0,390.0,901.0,403.0,1.4103,103900.0,NEAR BAY
-122.26,37.83,52.0,1470.0,330.0,689.0,309.0,3.48,191400.0,NEAR BAY
-122.26,37.83,52.0,2432.0,715.0,1377.0,696.0,2.5898,176000.0,NEAR BAY
-122.26,37.83,52.0,1665.0,419.0,946.0,395.0,2.0978,155400.0,NEAR BAY
-122.26,37.83,51.0,936.0,311.0,517.0,249.0,1.2852,150000.0,NEAR BAY
-122.26,37.84,49.0,713.0,202.0,462.0,189.0,1.025,118800.0,NEAR BAY
-122.26,37.84,52.0,950.0,202.0,467.0,198.0,3.9643,188800.0,NEAR BAY
-122.26,37.83,52.0,1443.0,311.0,660.0,292.0,3.0125,184400.0,NEAR BAY
-122.26,37.83,52.0,1656.0,420.0,718.0,382.0,2.6768,182300.0,NEAR BAY
-122.26,37.83,50.0,1125.0,322.0,616.0,304.0,2.026,142500.0,NEAR BAY
-122.27,37.82,43.0,1007.0,312.0,558.0,253.0,1.7348,137500.0,NEAR BAY
-122.26,37.82,40.0,624.0,195.0,423.0,160.0,0.9506,187500.0,NEAR BAY
-122.27,37.82,40.0,946.0,375.0,700.0,352.0,1.775,112500.0,NEAR BAY

671+A



데이터 가져오기

```
In [2]: housing = pd.read_csv('./housing.csv')
housing.head()
```

Out [2]:


	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	median_income	median_house_value	ocean_proximity
0	-122.23	37.88	52.0	1410	361	32082	15.3654	452600.0	NEAR BAY
1	-122.22	37.86	50.0	1080	200	32088	17.2574	358500.0	NEAR BAY
2	-122.24	37.85	50.0	1410	361	32088	17.2574	352100.0	NEAR BAY
3	-122.25	37.85	52.0	1080	200	32088	5.6431	341300.0	NEAR BAY
4	-122.25	37.85	52.0	1080	200	32088	3.8462	342200.0	NEAR BAY

pandas 라이브러리의 `read_csv` 함수를 사용해 데이터를 가져왔습니다.

| 주어진 데이터의 상위 5개 행을 표시하는 코드

| `read_csv` 함수 내부에는 데이터 파일의 경로가 지정되어야 합니다.

데이터 가져오기



데이터 가져오기

```
In [4]: import os
import tarfile
from six.moves import urllib

Download_root = "https://raw.githubusercontent.com/ExcelisorCJH/Hands-On-ML/master/"
Housing_path = "datasets/housing"
Housing_url = Download_root + Housing_path + "/housing.tgz"

def fetch_housing_data(housing_url = Housing_url, housing_path = Housing_path):
    if not os.path.isdir(housing_path):
        os.makedirs(housing_path)
    tgz_path = os.path.join(housing_path, "housing.tgz")
    urllib.request.urlretrieve(housing_url, tgz_path)
    housing_tgz = tarfile.open(tgz_path)
    housing_tgz.extractall(path=housing_path)
    housing_tgz.close()
```

```
In [5]: def load_housing_data(housing_path = Housing_path):
    csv_path = os.path.join(housing_path, "housing.csv")
    return pd.read_csv(csv_path)
```

```
In [6]: fetch_housing_data()
housing = load_housing_data()
housing.head()
```

데이터 가져오기

```
In [6]: fetch_housing_data()  
housing = load_housing_data()  
housing.head()
```

Out [6]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	ocean_proximity
0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0	8.3252	452600.0	NEAR BAY
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0	8.3014	358500.0	NEAR BAY
2	-122.24	37.85	52.0	1467.0	190.0	496.0	177.0	7.2574	352100.0	NEAR BAY
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0	5.6431	341300.0	NEAR BAY
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0	3.8462	342200.0	NEAR BAY

pandas 라이브러리의 `read_csv` 함수를 활용하여 csv파일을 불러왔습니다.

| 주어진 데이터의 상위 5개 Attribute를 보여주는 코드 : `dataframe.head()`

| `read_csv` 함수 내부에는 불러올 csv파일의 경로가 지정되어야 합니다.

2

—

데이터 통찰/시각화

데이터 시각화

```
In [7]: display(housing.info())
display(housing.describe())
```










```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
 0   longitude        20640 non-null   float64
 1   latitude         20640 non-null   float64
 2   housing_median_age 20640 non-null   float64
 3   total_rooms      20640 non-null   float64
 4   total_bedrooms   20433 non-null   float64
 5   population       20640 non-null   float64
 6   households       20640 non-null   float64
 7   median_income    20640 non-null   float64
 8   median_house_value 20640 non-null   float64
 9   ocean_proximity  20640 non-null   object  
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

```
None
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value
count	20640.000000	20640.000000	20640.000000	20640.000000	20433.000000	20640.000000	20640.000000	20640.000000	20640.000000
mean	-119.569704	35.631861	28.639486	2635.763081	537.870553	1425.476744	499.539680	3.870671	206855.816909
std	2.003532	2.135952	12.585558	2181.615252	421.385070	1132.462122	382.329753	1.899822	115395.615874
min	-124.350000	32.540000	1.000000	2.000000	1.000000	3.000000	1.000000	0.499900	14999.000000
25%	-121.800000	33.930000	18.000000	1447.750000	296.000000	787.000000	280.000000	2.563400	119600.000000
50%	-118.490000	34.260000	29.000000	2127.000000	435.000000	1166.000000	409.000000	3.534800	179700.000000
75%	-118.010000	37.710000	37.000000	3148.000000	647.000000	1725.000000	605.000000	4.743250	264725.000000
max	-114.310000	41.950000	52.000000	39320.000000	6445.000000	35682.000000	6082.000000	15.000100	500001.000000


데이터 시각화

```
In [8]: housing.hist(bins = 50, figsize = (20,15))  
plt.show()
```



데이터 시각화

히스토그램 정리



| 중간 소득 특성이 US달러가 아님!

→ 중간 소득이 최대 15달러일수는 없다!


| (중요!) 이미 Scaling 된 결과라는 판단이 필요합니다.

최소값을 0, 최대값을 15로 Scaling을 했네요

| 어느정도 정규성을 보장하고 있는 데이터라는 것을 이해합시다.

머신 러닝에서 전처리된 데이터를 다루는 경우가 흔하고, 이것이 문제가 되지는 않습니다. 다만, 데이터가 어떻게 계산되었는지는 반드시 이해합시다.


데이터 시각화



히스토그램 정리

| 중간 주택 연도/중간주택가격은 최대값, 최솟값을 한정

| (중요!) '중간주택가격'은 타깃 속성!




| 머신 러닝 학습 시 '50만 달러 초과 값'에 대한 예측은 나쁠 것

- 한계값 밖 구역에 대한 정확한 값을 찾는다 (불가능)
- 훈련 세트에서 이런 구역을 제거 (가능)

데이터 시각화

히스토그램 정리



| 전체적으로 첨도와 왜도가 상이합니다.


→ 이런 분포의 차이는 머신러닝 알고리즘에서 패턴을 찾기 어렵게 합니다.
→ 정규화/표준화가 필요합니다.

| 특성들의 스케일이 많이 다릅니다.


→ 각각의 서로 다른 특성들을 동일하게 Scaling할 필요가 있습니다.

데이터 시각화

```
In [10]: ax = housing.plot(kind='scatter', x='longitude', y='latitude')
ax.set(xlabel='경도', ylabel='위도');
```



```
In [11]: ax = housing.plot(kind='scatter', x='longitude', y='latitude',
ax.set(xlabel='경도', ylabel='위도');
```



데이터 시각화


	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	ocean_proximity	
0	-122.23	37.88		41.0	880.0	129.0	322.0	126.0	8.3252	452600.0	NEAR BAY
1	-122.22	37.86		21.0	7099.0	1106.0	2401.0	1138.0	8.3014	358500.0	NEAR BAY
2	-122.24	37.85		52.0	1467.0	190.0	496.0	177.0	7.2574	352100.0	NEAR BAY
3	-122.25	37.85		52.0	1274.0	235.0	558.0	219.0	5.6431	341300.0	NEAR BAY
4	-122.25	37.85		52.0	1627.0	280.0	565.0	259.0	3.8462	342200.0	NEAR BAY



```
In [13]: ax = housing.plot(kind='scatter', x='longitude', y='latitude', alpha=0.4,
s = housing['population']/100, label='인구', figsize=(10,7),
c = 'median_house_value', cmap=plt.get_cmap("jet"), colorbar=True, sharex=False)

ax.set(xlabel='경도', ylabel='위도')
plt.legend()

Out[13]: <matplotlib.legend.Legend at 0x1d60a300bc8>
```





데이터 통찰


```
In [13]: ax = housing.plot(kind='scatter', x='longitude', y='latitude', alpha=0.4,
s = housing['population']/100, label='인구', figsize=(10,7),
c = 'median_house_value', cmap=plt.get_cmap("jet"), colorbar=True, sharex=False)

ax.set(xlabel='경도', ylabel='위도')
plt.legend()
```

```
Out[13]: <matplotlib.legend.Legend at 0x1d60a300bc0>
```



데이터 통찰



데이터 통찰

```
In [14]: corr_matrix = housing.corr()  
corr_matrix
```

Out [14]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value
longitude	1.000000	-0.924664	-0.108197	0.044568	0.069608	0.099773	0.055310	-0.015176	-0.045967
latitude	-0.924664	1.000000	0.011173	-0.036100	-0.066983	-0.108785	-0.071035	-0.079809	-0.144160
housing_median_age	-0.108197	0.011173	1.000000	-0.361262	-0.320451	-0.296244	-0.302916	-0.119034	0.105623
total_rooms	0.044568	-0.036100	-0.361262	1.000000	0.930380	0.857126	0.918484	0.198050	0.134153
total_bedrooms	0.069608	-0.066983	-0.320451	0.930380	1.000000	0.877747	0.979728	-0.007723	0.049686
population	0.099773	-0.108785	-0.296244	0.857126	0.877747	1.000000	0.907222	0.004834	-0.024650
households	0.055310	-0.071035	-0.302916	0.918484	0.979728	0.907222	1.000000	0.013033	0.065843
median_income	-0.015176	-0.079809	-0.119034	0.198050	-0.007723	0.004834	0.013033	1.000000	0.688075
median_house_value	-0.045967	-0.144160	0.105623	0.134153	0.049686	-0.024650	0.065843	0.688075	1.000000


데이터 통찰

```
corr_matrix["median_house_value"].sort_values(ascending=False)
```

```
median_house_value      1.000000
median_income          0.688075
total_rooms            0.134153
housing_median_age    0.105623
households             0.065843
total_bedrooms         0.049686
population             -0.024650
longitude              -0.045967
latitude               -0.144160
Name: median_house_value, dtype: float64
```


데이터 통찰

```
In [16]: from pandas.plotting import scatter_matrix  
  
attributes = ["median_house_value", "median_income", "total_rooms",  
"housing_median_age"]  
  
scatter_matrix(housing[attributes], figsize=(12,8));
```




데이터 통찰

```
In [17]: housing.plot(kind="scatter", x="median_income", y="median_house_value",
                     alpha=0.1)
plt.axis([0, 16, 0, 550000]);
```



데이터 통찰

```
In [17]: housing.plot(kind="scatter", x="median_income", y="median_house_value",
                     alpha=0.1)
plt.axis([0, 16, 0, 550000]);
```



데이터 통찰

```
In [18]: housing['rooms_per_household'] = housing['total_rooms'] / housing['households']
housing['bedrooms_per_room'] = housing['total_bedrooms'] / housing['total_rooms']
housing['population_per_household'] = housing['population'] / housing['households']

corr_matrix = housing.corr()
corr_matrix["median_house_value"].sort_values(ascending=False)
```


```
Out [18]: median_house_value      1.000000
median_income                  0.688075
rooms_per_household           0.151948
total_rooms                     0.134153
housing_median_age            0.105623
households                     0.065843
total_bedrooms                  0.049686
population_per_household     -0.023737
population                      -0.024650
longitude                      -0.045967
latitude                       -0.144160
bedrooms_per_room              -0.255880
Name: median_house_value, dtype: float64
```

3


—

데이터 훈련/평가


데이터 훈련



데이터 훈련



데이터 훈련



데이터 훈련

데이터를 더 깊게 들여다 보기 전에 Test Set을 따로 떼어놓아야 합니다.

↳ 언더피팅 (Underfitting)의 위험 존재

그리고 Test Set을 절대 들여다보면 안됩니다!

| WHY?

- 테스트 셋에서 겉으로 드러난 어떤 패턴에 속는다.
- 특정 머신러닝 모델만 낙관적 편향에 따라 선택.
- 일반화했을 때, 기대한 성능이 나오지 않을 것 (데이터 스누핑 발생).

데이터 훈련

Q

올바른 Training Set이란 무엇일까요?

어떻게 올바른 Training Set을 만들죠?

데이터 훈련 - 전처리

▶ 계층적 샘플링

```
In [19]: from sklearn.model_selection import train_test_split  
  
train_set, test_set = train_test_split(housing, test_size=0.2, random_state=42)  
print(len(train_set), "train +", len(test_set), "test")  
  
16512 train + 4128 test
```

```
In [18]: test_set
```

Out[18]:

	longitude	latitude	housing_median_age	total_rooms	...	ocean_proximity	rooms_per_household	bedrooms_per_room	population_per_household
20046	-119.01	36.06	25.0	1505.0	...	INLAND	4.1...	NaN	3.8...
3024	-119.46	35.14	30.0	2943.0	...	INLAND	5.0...	NaN	2.6...
15663	-122.44	37.80	52.0	3830.0	...	NEA...	3.9...	NaN	1.3...
20484	-118.72	34.28	17.0	3051.0	...	<1H...	6.1...	NaN	3.4...
9814	-121.93	36.62	34.0	2351.0	...	NEA...	5.4...	NaN	2.4...
...
15362	-117.22	33.36	16.0	3165.0	...	<1H...	7.0...	0.1...	2.9...
16623	-120.83	35.36	28.0	4323.0	...	NEA...	6.1...	0.2...	2.3...
18086	-122.05	37.31	25.0	4111.0	...	<1H...	7.2...	0.1...	2.7...
2144	-119.76	36.77	36.0	2507.0	...	INLAND	5.2...	0.1...	2.5...
3665	-118.37	34.22	17.0	1787.0	...	<1H...	3.9...	0.2...	3.7...

4128 rows × 13 columns

데이터 훈련 - 전처리

▶ 계층적 샘플링

```
In [19]: train_set
```

```
Out[19]:
```


	longitude	latitude	housing_median_age	total_rooms	...	ocean_proximity	rooms_per_household	bedrooms_per_room	population_per_household	
14196	-117.03	32.71	33.0	3126.0	...	NEA...	5.0...	0.2...	3.6...	
8267	-118.16	33.77	49.0	3382.0	...	NEA...	4.4...	0.2...	1.7...	
17445	-120.48	34.66	4.0	1897.0	...	NEA...	5.6...	0.1...	2.7...	
14265	-117.11	32.69	36.0	1421.0	...	NEA...	4.0...	0.2...	3.9...	
2271	-119.80	36.78	43.0	2382.0	...	INLAND	6.2...	0.1...	2.3...	
...
11284	-117.96	33.78	35.0	1330.0	...	<1H...	6.1...	0.1...	3.0...	
11964	-117.43	34.02	33.0	3084.0	...	INLAND	6.8...	0.1...	3.9...	
5390	-118.38	34.03	36.0	2101.0	...	<1H...	3.9...	0.2...	3.3...	
860	-121.96	37.58	15.0	3575.0	...	<1H...	6.3...	0.1...	3.1...	
15795	-122.42	37.77	52.0	4226.0	...	NEA...	3.4...	0.3...	2.1...	

16512 rows × 13 columns

데이터 훈련 - 전처리

▶ 계층적 샘플링

```
In [20]: housing['median_income'].hist();
```



히스토그램 정리

- | 정규성을 띠는 데이터 특성인 중간 소득을 살펴봅시다!
 - 소득 별 빈도가 상이합니다.
 - 이러한 상황에서 훈련 셋 80%, 테스트 셋 20%를 **아무 렇게나 뽑는 것은 옳지 않습니다.**
- | 계층별로 트레이닝, 테스트 셋을 뽑도록 합시다!

데이터 훈련 - 전처리


▶ 계층적 샘플링

```
In [23]: housing['income_cat'] = np.ceil(housing['median_income'] / 1.5)
housing['income_cat'].where(housing['income_cat'] < 5, 5.0, inplace=True)
```

```
In [24]: housing['income_cat'].value_counts()
```

```
Out[24]: 3.0    7236
2.0    6581
4.0    3639
5.0    2362
1.0    822
Name: income_cat, dtype: int64
```

```
In [25]: housing['income_cat'].hist();
```



히스토그램 정리

| 데이터를 5계층으로 분류했습니다.

→ 계층 별 비중에 따라 Sampling을 진행합시다!

데이터 훈련 - 전처리

▶ 계층적 샘플링

비슷한 비중으로 Sampling 되었습니다!

```
In [29]: from sklearn.model_selection import StratifiedShuffleSplit
split = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=42)
for train_index, test_index in split.split(housing, housing['income_cat']):
    strat_train_set = housing.loc[train_index]
    strat_test_set = housing.loc[test_index]
```

```
display(housing['income_cat'].value_counts()/len(housing))
print('='*20)
display(strat_test_set['income_cat'].value_counts()/len(strat_test_set))
```

```
3.0    0.350581
2.0    0.318847
4.0    0.176308
5.0    0.114438
1.0    0.039826
Name: income_cat, dtype: float64
=====

```

```
3.0    0.350533
2.0    0.318798
4.0    0.176357
5.0    0.114583
1.0    0.039729
Name: income_cat, dtype: float64
```

데이터 훈련 - 전처리

▶ 계층적 샘플링

```
In [32]: def income_cat_proportions(data):
    return data["income_cat"].value_counts() / len(data)

train_set, test_set = train_test_split(housing, test_size = 0.2, random_state=42)

compare_props = pd.DataFrame({
    "Overall" : income_cat_proportions(housing),
    "Stratified": income_cat_proportions(strat_test_set),
    "Random" : income_cat_proportions(test_set),
}).sort_index()

compare_props['Rand. %error'] = 100 * compare_props["Random"] / compare_props["Overall"] - 100
compare_props["Strat. %error"] = 100 * compare_props["Stratified"] / compare_props["Overall"] - 100

compare_props
```

Out [32]:

	Overall	Stratified	Random	Rand. %error	Strat. %error
1.0	0.039826	0.039729	0.040213	0.973236	-0.243309
2.0	0.318847	0.318798	0.324370	1.732260	-0.015195
3.0	0.350581	0.350533	0.358527	2.266446	-0.013820
4.0	0.176308	0.176357	0.167393	-5.056334	0.027480
5.0	0.114438	0.114583	0.109496	-4.318374	0.127011

테이블 정리

| 계층 별 비중의 오차 Table입니다.

→ Rand. %error > Strat. %error

→ 계층화하고 데이터를 선택하는 것의 일 반화 성능이 더 좋을 것입니다.

데이터 훈련 - 전처리

▶ 계층적 샘플링

```
In [35]: for set_ in (strat_train_set, strat_test_set, housing):
    set_.drop("income_cat", axis=1, inplace=True)

housing.columns
```

```
Out[35]: Index(['longitude', 'latitude', 'housing_median_age', 'total_rooms',
       'total_bedrooms', 'population', 'households', 'median_income',
       'median_house_value', 'ocean_proximity', 'rooms_per_household',
       'bedrooms_per_room', 'population_per_household'],
      dtype='object')
```

| 더 이상 쓸모가 없는 'Income_cat' 특성을 제거합시다.

데이터 훈련 - 전처리

▶ NULL값 제거

```
In [57]: housing = strat_train_set.drop(columns=['median_house_value']) #train_X  
housing_labels = strat_train_set['median_house_value'].copy() #train_Y
```

```
In [59]: housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 16512 entries, 17606 to 15775  
Data columns (total 12 columns):  
 #   Column           Non-Null Count  Dtype     
 --  --  
 0   longitude        16512 non-null  float64  
 1   latitude         16512 non-null  float64  
 2   housing_median_age 16512 non-null  float64  
 3   total_rooms      16512 non-null  float64  
 4   total_bedrooms   16354 non-null  float64  
 5   population       16512 non-null  float64  
 6   households       16512 non-null  float64  
 7   median_income    16512 non-null  float64  
 8   ocean_proximity  16512 non-null  object    
 9   rooms_per_household 16512 non-null  float64  
 10  bedrooms_per_room 16354 non-null  float64  
 11  population_per_household 16512 non-null  float64  
dtypes: float64(11), object(1)  
memory usage: 1.6+ MB
```

데이터 훈련 - 전처리

▶ NULL값 제거

```
In [74]: housing[housing['bedrooms_per_room'].isna() == True]
```

```
Out[74]:
```

	longitude	latitude	housing_median_age	total_rooms	...	ocean_proximity	rooms_per_household	bedrooms_per_room	population_per_household
4629	-118.30	34.07	18.0	3759.0	...	<1H...	2.5...	NaN	2.2...
6068	-117.86	34.01	16.0	4632.0	...	<1H...	6.3...	NaN	4.1...
17923	-121.97	37.35	30.0	1955.0	...	<1H...	5.0...	NaN	2.5...
13656	-117.30	34.05	6.0	2155.0	...	INLAND	5.5...	NaN	2.6...
19252	-122.79	38.48	7.0	6837.0	...	<1H...	4.8...	NaN	2.4...
...
3376	-118.28	34.25	29.0	2559.0	...	<1H...	3.3...	NaN	2.4...
4691	-118.37	34.07	50.0	2519.0	...	<1H...	4.8...	NaN	2.1...
6052	-117.76	34.04	34.0	1914.0	...	INLAND	5.8...	NaN	4.7...
17198	-119.75	34.45	6.0	2864.0	...	NEA...	4.7...	NaN	2.3...
4738	-118.38	34.05	49.0	702.0	...	<1H...	3.7...	NaN	2.4...

158 rows × 12 columns

제거? 대치?

데이터 훈련 - 전처리

▶ NULL값 제거

```
In [76]: from sklearn.impute import SimpleImputer  
  
housing_num = housing.drop('ocean_proximity', axis=1)
```

```
In [77]: imputer = SimpleImputer(strategy='median')  
imputer.fit(housing_num)
```

```
Out[77]: SimpleImputer(strategy='median')
```

```
In [81]: print('imputer.statistics_:', list(map(lambda x: "%.<2f"%x, imputer.statistics_)))  
print('housing_num.median().values:', list(map(lambda x: "%.<2f"%x, housing_num.median().values)))
```

```
imputer.statistics_ : ['-118.51', '34.26', '29.00', '2119.50', '433.00', '1164.00', '408.00', '3.54', '5.23', '0.20', '2.82']  
housing_num.median().values : ['-118.51', '34.26', '29.00', '2119.50', '433.00', '1164.00', '408.00', '3.54', '5.23', '0.20', '2.82']
```

데이터 훈련 - 전처리

▶ NULL값 제거




```
In [83]: X = imputer.transform(housing_num)  
housing_tr = pd.DataFrame(X, columns = housing_num.columns, index = housing.index.values)
```

```
In [84]: housing_tr.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 16512 entries, 17606 to 15775  
Data columns (total 11 columns):  
 #   Column           Non-Null Count  Dtype     
 ---  --  
 0   longitude        16512 non-null   float64  
 1   latitude         16512 non-null   float64  
 2   housing_median_age 16512 non-null   float64  
 3   total_rooms      16512 non-null   float64  
 4   total_bedrooms   16512 non-null   float64  
 5   population       16512 non-null   float64  
 6   households       16512 non-null   float64  
 7   median_income    16512 non-null   float64  
 8   rooms_per_household 16512 non-null   float64  
 9   bedrooms_per_room 16512 non-null   float64  
 10  population_per_household 16512 non-null   float64  
dtypes: float64(11)  
memory usage: 1.5 MB
```

데이터 훈련 - 전처리


▶ 이상치 제거




데이터 훈련 - 전처리

▶ 이상치 제거

```
In [63]: housing_labels.hist(bins = 100, figsize = (20,15))  
plt.show()
```



```
In [70]: housing_tr['housing_median_age'].hist(bins = 50, figsize = (20,15))  
plt.show()
```



데이터 훈련 - 전처리

▶ 이상치 제거

```
In [83]: housing_tr[housing_tr['housing_median_age'] == max(housing_tr['housing_median_age'])]
```

```
Out[83]:
```

	longitude	latitude	housing_median_age	total_rooms	...	median_income	rooms_per_household	bedrooms_per_room	population_per_household
8879	-118.50	34.04	52.0	2233.0	...	8.3839	8.0...	0.1...	2.7...
16121	-122.46	37.79	52.0	2059.0	...	3.7419	5.1...	0.2...	2.4...
2651	-124.10	40.47	52.0	1196.0	...	3.5345	4.5...	0.1...	3.6...
16131	-122.48	37.79	52.0	4683.0	...	4.1148	4.8...	0.2...	2.3...
533	-122.27	37.78	52.0	1222.0	...	3.7708	4.6...	0.2...	2.3...
...
15670	-122.44	37.80	52.0	1006.0	...	2.7717	3.9...	0.2...	1.7...
11878	-117.38	33.99	52.0	1797.0	...	2.7054	5.7...	0.1...	2.8...
108	-122.24	37.82	52.0	3481.0	...	3.9000	4.8...	0.2...	2.0...
5364	-118.42	34.04	52.0	1358.0	...	5.6454	5.0...	0.2...	2.1...
15775	-122.45	37.77	52.0	3095.0	...	3.5750	4.8...	0.2...	1.9...

1027 rows × 11 columns

데이터 훈련 - 전처리

▶ 이상치 제거

```
In [84]: housing_labels[housing_labels.values==max(housing_labels.values)]
```

```
Out [84]: 8879      500...
 4861      500...
 15614     500...
 18086     500...
 16121     500...
 ...
 8667      500...
 5306      500...
 18067     500...
 5364      500...
 15775     500...
Name: median_house_value, Length: 786, dtype: float64
```

데이터 훈련 - 전처리

▶ 이상치 제거

```
In [82]: X_index = housing_tr[housing_tr['housing_median_age'] == max(housing_tr['housing_median_age'])].index  
y_index = housing_labels[housing_labels.values == max(housing_labels.values)].index  
print("Variable_index - Target_index : ", len(set(X_index) - set(y_index)))  
print("Target_index - Variable_index : ", len(set(y_index) - set(X_index)))  
print("Union of Variable_index & Target_index : ", len(set(y_index).union(set(X_index))))
```

```
Variable_index - Target_index : 887  
Target_index - Variable_index : 646  
Union of Variable_index & Target_index : 1673
```

데이터 훈련 - 전처리


▶ 이상치 제거

```
In [92]: tmp_x = housing_tr.drop(index = list(X_union_y))
tmp_y = housing_labels.drop(index = list(X_union_y))


tmp_x.hist(bins = 50, figsize = (20,15))
plt.show()

tmp_y.hist(bins = 50, figsize = (20,15))
plt.show()
```

X - variables




y - target



데이터 훈련 - 전처리

▶ 특성 스케일링



데이터 훈련 - 전처리

▶ 특성 스케일링

정규화

```
In [102]: from sklearn.preprocessing import MinMaxScaler  
  
data = [[-1, 2], [-0.5, 6], [0, 10], [1, 18]]  
scaler = MinMaxScaler()  
print(scaler.fit(data))  
print("max of data = ", scaler.data_max_ )  
print("min of data = ", scaler.data_min_ )  
print("data range = ", scaler.data_range_ )  
print("<data scaling> \n", scaler.transform(data))  
print("[2, 2] minmax scaling ", scaler.transform([[2,2]]))
```

```
MinMaxScaler()  
max of data =  [ 1. 18.]  
min of data =  [-1. 2.]  
data range =  [ 2. 16.]  
<data scaling>  
[[0. 0.]  
[0.25 0.25]  
[0.5 0.5 ]  
[1. 1. ]]  
[2, 2] minmax scaling  [[1.5 0. ]]
```

표준화

```
In [104]: from sklearn.preprocessing import StandardScaler  
  
data = [[-1, 2], [-0.5, 6], [0, 10], [1, 18]]  
scaler = StandardScaler()  
print(scaler.fit(data))  
print("mean = ", scaler.mean_ )  
print("variance = ", scaler.var_ )  
print("z-score \n", scaler.transform(data))  
print("[2, 2] standard scaling ", scaler.transform([[2,2]]))
```

```
StandardScaler()  
mean =  [-0.125  9. ]  
variance =  [ 0.546875 35. ]  
z-score  
[[ -1.18321596 -1.18321596]  
[ -0.50709255 -0.50709255]  
[  0.16903085  0.16903085]  
[  1.52127766  1.52127766]]  
[2, 2] standard scaling  [[ 2.87352447 -1.18321596]]
```

데이터 훈련 - 전처리

▶ 특성 스케일링

```
In [58]: scaler = StandardScaler()
scaler.fit(tmp_x)
tmp_prepared = scaler.transform(tmp_x)

prepared = pd.DataFrame(tmp_prepared, columns = tmp_x.columns, index = tmp_x.index)
prepared.head()
```

Out [58]:

	longitude	latitude	housing_median_age	total_rooms	...	median_income	rooms_per_household	bedrooms_per_room	population_per_household	
17606	-1.2...	0.7...		0.9...	-0.4...	...	-0.6...	-0.3...	0.1...	-0.0...
18632	-1.2...	0.6...		-1.1...	-0.9...	...	1.7...	0.2...	-1.0...	-0.0...
14650	1.1...	-1.3...		0.3...	-0.3...	...	-0.5...	-0.4...	0.5...	-0.0...
3230	-0.0...	0.3...		-0.1...	-0.3...	...	-1.1...	-0.0...	-0.2...	0.0...
3555	0.4...	-0.6...		-0.8...	1.8...	...	-0.4...	-0.3...	0.3...	-0.0...

5 rows × 11 columns

데이터 훈련 - 전처리

▶ 범주형 데이터 처리

```
In [55]: housing['ocean_proximity']
```

```
Out[55]:
```

17606	<1H...
18632	<1H...
14650	NEA...
3230	INLAND
3555	<1H...
	...
6563	INLAND
12053	INLAND
13908	INLAND
11159	<1H...
15775	NEA...

Name: ocean_proximity, Length: 16512, dtype: object

```
In [54]: housing_cat = pd.get_dummies(housing['ocean_proximity'])  
housing_cat = housing_cat.loc[tmp_x.index]  
housing_cat
```

```
Out[54]:
```

	<1H OCEAN	INLAND	ISLAND	NEAR BAY	NEAR OCEAN
17606	1	0	0	0	0
18632	1	0	0	0	0
14650	0	0	0	0	1
3230	0	1	0	0	0
3555	1	0	0	0	0
...
7364	1	0	0	0	0
6563	0	1	0	0	0
12053	0	1	0	0	0
13908	0	1	0	0	0
11159	1	0	0	0	0

14839 rows × 5 columns

데이터 훈련 - 전처리

▶ 범주형 데이터 처리

```
In [57]: housing = pd.concat([prepared, housing_cat], axis=1)  
housing
```

Out [57]:

	longitude	latitude	housing_median_age	total_rooms	...	INLAND	ISLAND	NEAR BAY	NEAR OCEAN
17606	-1.2...	0.7...	0.9...	-0.4...	...	0	0	0	0
18632	-1.2...	0.6...	-1.1...	-0.9...	...	0	0	0	0
14650	1.1...	-1.3...	0.3...	-0.3...	...	0	0	0	1
3230	-0.0...	0.3...	-0.1...	-0.3...	...	1	0	0	0
3555	0.4...	-0.6...	-0.8...	1.8...	...	0	0	0	0
...
7364	0.6...	-0.7...	1.4...	-0.8...	...	0	0	0	0
6563	0.6...	-0.6...	1.6...	-0.6...	...	1	0	0	0
12053	0.9...	-0.7...	1.1...	-0.6...	...	1	0	0	0
13908	1.5...	-0.7...	-1.5...	1.0...	...	1	0	0	0
11159	0.7...	-0.8...	0.3...	-0.3...	...	0	0	0	0

14839 rows × 16 columns

데이터 훈련 - 훈련

▶ 선형 회귀 모델

```
In [62]: from sklearn.linear_model import LinearRegression  
  
housing_prepared = np.array(housing)  
housing_labels = tmp_y  
  
lin_reg = LinearRegression()  
lin_reg.fit(housing_prepared, housing_labels)
```

```
Out[62]: LinearRegression()
```

```
In [64]: from sklearn.metrics import mean_squared_error  
  
housing_predictions = lin_reg.predict(housing_prepared)  
lin_mse = mean_squared_error(housing_labels, housing_predictions)  
lin_rmse = np.sqrt(lin_mse)  
print('RMSE : ', lin_rmse)
```

```
RMSE : 58244.36758810236
```

데이터 훈련 - 훈련

▶ 의사결정나무 모델

```
In [66]: from sklearn.tree import DecisionTreeRegressor  
  
tree_reg = DecisionTreeRegressor()  
tree_reg.fit(housing_prepared, housing_labels)
```


```
Out[66]: DecisionTreeRegressor()
```

```
In [67]: housing_predictions = tree_reg.predict(housing_prepared)  
tree_mse = mean_squared_error(housing_labels, housing_predictions)  
tree_rmse = np.sqrt(tree_mse)  
print('RMSE : ', tree_rmse)
```

```
RMSE :  0.0
```

데이터 훈련 - 훈련

▶ Cross Validation



데이터 훈련 - 훈련

▶ Cross Validation

```
In [69]: from sklearn.model_selection import cross_val_score

scores = cross_val_score(tree_reg, housing_prepared, housing_labels, scoring='neg_mean_squared_error', cv=10)
tree_rmse_scores = np.sqrt(-scores)
```

```
In [70]: def display_scores(scores):
    print("점수 :", scores)
    print("평균 :", scores.mean())
    print("표준편차 :", scores.std())

display_scores(tree_rmse_scores)
```

```
점수 : [62454.41054262 63025.77955203 63832.92408127 63777.66473357
65482.47275832 65252.96917991 60402.82013775 61480.54712836
64392.03096883 65620.55714711]
평균 : 63572.21762297725
표준편차 : 1655.9051509908227
```

CrossValidation - Tree model

```
In [71]: lin_scores = cross_val_score(lin_reg, housing_prepared, housing_labels, scoring='neg_mean_squared_error', cv=10)

lin_rmse_scores = np.sqrt(-lin_scores)
display_scores(lin_rmse_scores)
```

```
점수 : [56722.17148984 58053.7705809 57735.81527587 64063.65140457
58965.20314997 59435.21574818 54764.97078541 58147.54954935
60177.86191379 57146.94472927]
평균 : 58521.31546271383
표준편차 : 2338.912842506291
```

CrossValidation - Linear Regression model

데이터 훈련 - 훈련

▶ Cross Validation


```
In [70]: def display_scores(scores):
    print("점수 :", scores)
    print("평균 :", scores.mean())
    print("표준편차 :", scores.std())
```

```
display_scores(tree_rmse_scores)
```

```
점수 : [62454.41054262 63025.77955203 63832.92408127 63777.66473357
65482.47275832 65252.96917991 60402.82013775 61480.54712836
64392.03096883 65620.55714711]
평균 : 63572.21762297725
표준편차 : 1655.9051509908227
```

```
In [71]: lin_scores = cross_val_score(lin_reg, housing_prepared, housing_labels, scoring='neg_mean_squared_error', cv=10)
lin_rmse_scores = np.sqrt(-lin_scores)
display_scores(lin_rmse_scores)
```

```
점수 : [56722.17148984 58053.7705809 57735.81527587 64063.65140457
58965.20314997 59435.21574818 54764.97078541 58147.54954935
60177.86191379 57146.94472927]
평균 : 58521.31546271383
표준편차 : 2338.912842506291
```



어떤 알고리즘이 더 좋을까요?

Test Set과의 비교

다음시간에...

Q&A